

NS-2/tcptrace Project Report

ECE 4607

1 May, 2015

Enmao Diao


James Eccles

Daniel Graves

Wes Smith

Section I - Tcl script located in Appendix A

1. The following table shows the number of TCP segments received at the destination and the overall throughput for each flow. Simulations performed for TCP New Reno, Reno, SACK, and Tahoe.



Network Topology

In our simulations, there are four flows: Flow 1 is S1 - D1, Flow 2 is S1 - D2, Flow 3 is S2 - D1, Flow 4 is S2 - D2.

Table I


TCP Segments and Overall Throughput



| *Throughput in Mbps | | Flow | | | |
|---------------------|------------|--------|--------|--------|--------|
| TCP Variant | | 1 | 2 | 3 | 4 |
| New Reno | Segments | 13351 | 45852 | 19096 | 27839 |
| | Throughput | 1.1312 | 3.8850 | 1.6179 | 2.3588 |
| Reno | Segments | 14625 | 38180 | 19544 | 30046 |
| | Throughput | 1.2390 | 3.2347 | 1.6558 | 2.5456 |
| SACK | Segments | 14477 | 38759 | 20207 | 27851 |


| | | | | | |
|-------|------------|--------|--------|--------|--------|
| | Throughput | 1.2266 | 3.2840 | 1.7121 | 2.3597 |
| Tahoe | Segments | 14477 | 38759 | 20207 | 27851 |
| | Throughput | 1.2266 | 3.2840 | 1.7121 | 2.3597 |

2. Plots of throughput and congestion window vs time of each flow:


New Reno











Reno











SACK










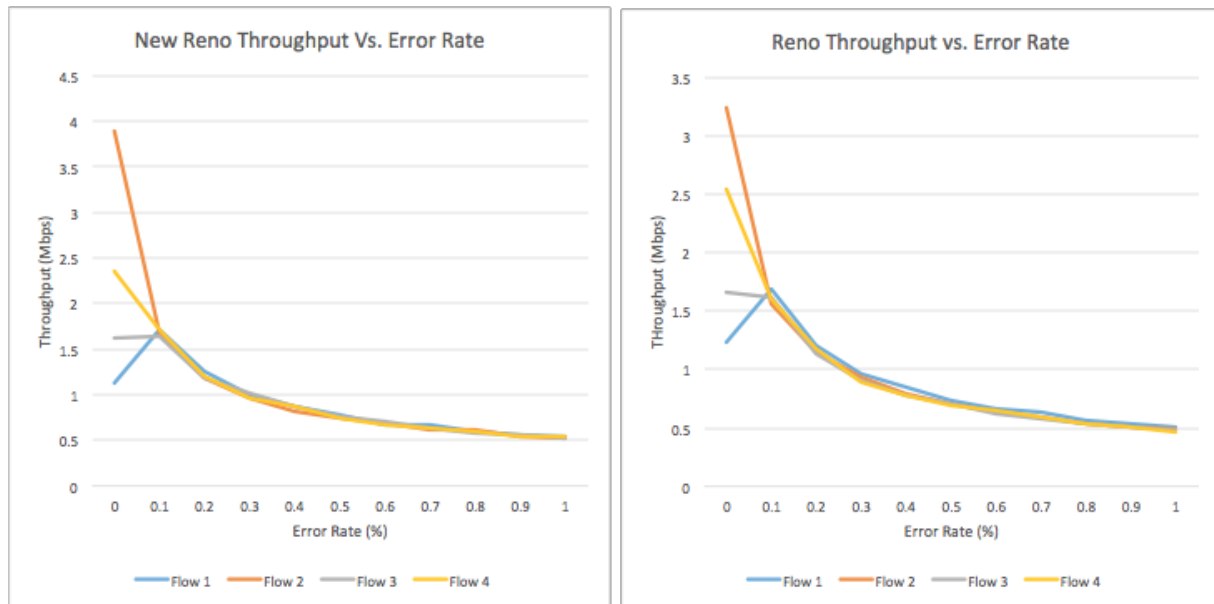
Tahoe

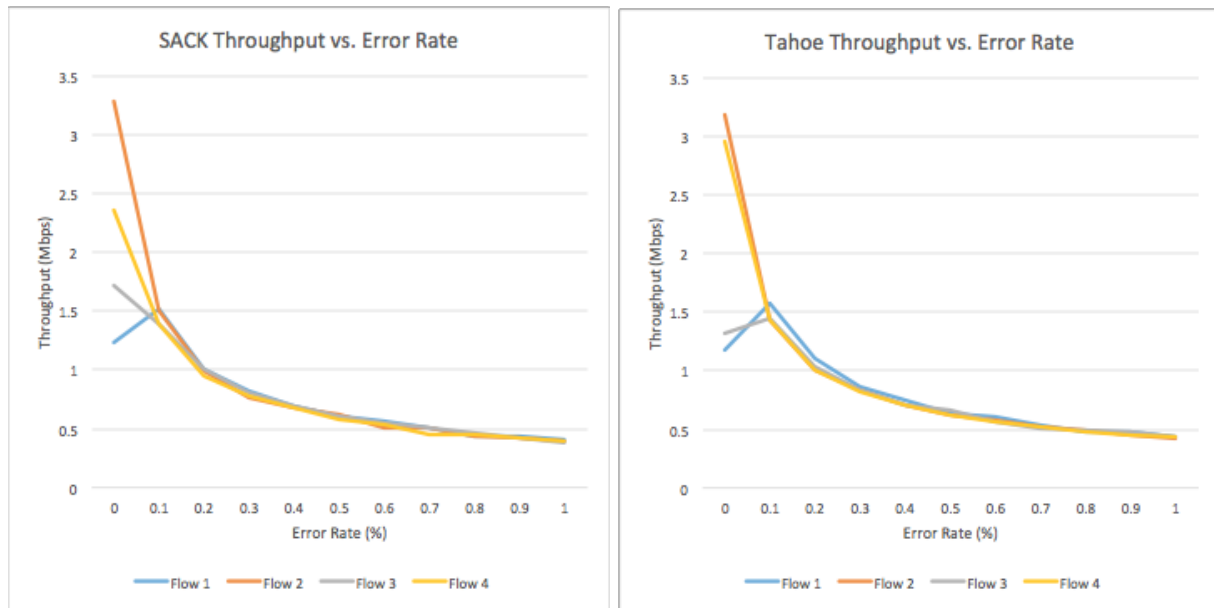






3. A loss module was added into the network between nodes R1 and R2. Average throughput was measured twenty times for each flow for error rates from 0% to 1%. Awk and Shell file are used and shown in Appendix C and D.





4. At zero error rate, Tahoe provides more throughput to Flow 4 between S2-D2 and less throughput to Flow 3 between S2-D1. Meanwhile, new Reno provides less throughput for Flow 4 between S2-D2 and adds more throughput to Flow 2 between S1-D2. As error rate increases beyond 0%, the four TCP flavors begin to perform similarly but with small differences in maximum throughput. At an error rate of 1%, the four New Reno flows each average 0.5289 Mbps. At this same error rate, each Reno flow averages 0.4844 Mbps, each Tahoe flow averages 0.4316 Mbps, and each SACK flow averages 0.3931 Mbps. This means that New Reno performs the best as error rates increase, and SACK performs the worst.

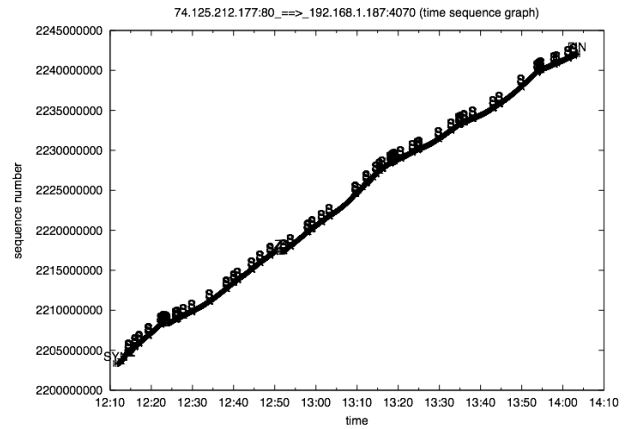
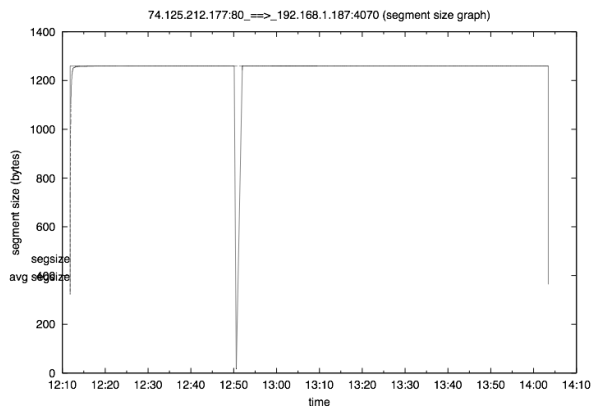

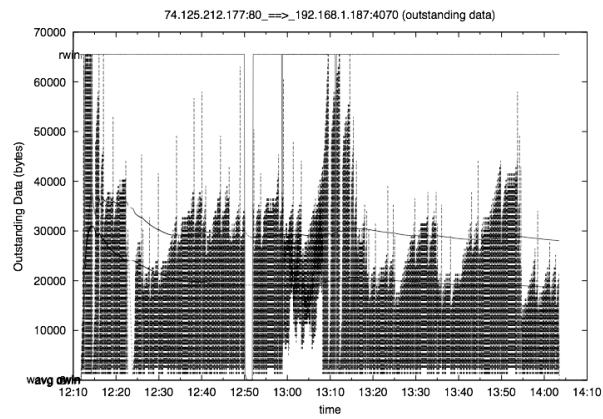
Tahoe suffers in lossy environments because it requires a full timeout interval to detect a packet loss. Reno adds a Fast Retransmit mode, which detects packet loss earlier by counting duplicate ACKs. New Reno is a slight modification of Reno, which may perform better in lossy environments by refraining from excessive congestion window

reductions. SACK may suffer in high-error environments from overusing channel resources on acknowledgments.

Section 2 - tcpdump output located in Appendix B

1. IP addresses involved:
 - a. Local: 192.169.1.187
 - b. Remote: 74.125.212.177
2. Port numbers of local and remote hosts:
 - a. Local: 4070
 - b. Remote: 80
3. Max Segment Size (MSS) used in either direction:
 - a. 1260 bytes
4. Total bytes and unique bytes sent by remote end-host:
 - a. Total (actual) bytes: 38785361
 - b. Unique bytes: 38719841
5. Number of unique data pkts sent by remote end-host:
 - a. 30785 pkts
6. Average downlink throughput (by-hand calculation):
 - a. $30790 \text{ pkts} / 111.603 \text{ sec} = 275.89 \text{ pkts} / \text{sec}$
7. Avg RTT in either direction:
 - a. Local->Remote: 30.8 ms
 - b. Remote->Local: 49.9 ms
8. Minimum end-to-end RTT:
 - a. 30.6 ms
9. TCP flavors used by both end-hosts. How do you identify them?
 - a. Local (a): **TCP SACK**. We know this because it uses sacks (see sack pkts sent in Appendix B).
 - b. Remote (b): RTT spikes at 12:23, 12:50, 13:12. Look at outstanding window at this point. After 12:20, outstanding data (estimates congestion window at sender) is very low (slow start?). Quickly returns to threshold. Same at 12:50. At 13:12, outstanding data drops and returns nearly immediately to its high window. This could indicate that the fast recovery algorithm is being used. 12:20 and 12:50

definitely indicate slow-start modes, therefore, the remote server uses **TCP Tahoe**.



10. Maximum congestion control window size used by remote end-host:

a. 8896 bytes

Appendix A

Tcl script for TCP New Reno (other TCP variants, change the Agent accordingly):

```
#Create a simulator object
set ns [new Simulator]

#Define different colors for data flows (for NAM)
$ns color 1 Blue
$ns color 2 Blue
$ns color 3 Red
$ns color 4 Red

#Open the NAM trace file
set nf [open projout.nam w]
$ns namtrace-all $nf

#Congestion Window file
set outfile1 [open "WinFile1" w]
set outfile2 [open "WinFile2" w]
set outfile3 [open "WinFile3" w]
set outfile4 [open "WinFile4" w]

#Define a 'finish' procedure
proc finish {} {
    global ns nf
    $ns flush-trace
    #Close the NAM trace file
    close $nf
    #Execute NAM on the trace file
    exec nam projout.nam &
    exit 0
}

#Create seven nodes
set n0 [$ns node]
set n1 [$ns node]
set n2 [$ns node]
set n3 [$ns node]
set n4 [$ns node]
set n5 [$ns node]
set n6 [$ns node]

#Create links between the nodes
$ns duplex-link $n0 $n2 20Mb 4ms DropTail
$ns duplex-link $n1 $n2 10Mb 5ms DropTail
$ns duplex-link $n2 $n3 20Mb 40ms DropTail
$ns duplex-link $n3 $n4 10Mb 45ms DropTail
$ns duplex-link $n4 $n5 20Mb 3ms DropTail
$ns duplex-link $n4 $n6 10Mb 5ms DropTail

#Set Queue Size of all links to 50
$ns queue-limit $n0 $n2 50
$ns queue-limit $n1 $n2 50
$ns queue-limit $n2 $n3 50
$ns queue-limit $n3 $n4 50
$ns queue-limit $n4 $n5 50
$ns queue-limit $n4 $n6 50

#Give node position (for NAM)
```

```
$ns duplex-link-op $n0 $n2 orient right-down
$ns duplex-link-op $n1 $n2 orient right-up
$ns duplex-link-op $n2 $n3 orient right
$ns duplex-link-op $n3 $n4 orient right
$ns duplex-link-op $n4 $n5 orient right-up
$ns duplex-link-op $n4 $n6 orient right-down
```

```
#Setup TCP connection S1-D1
set tcp0 [new Agent/TCP/Newreno]
$ns attach-agent $n0 $tcp0
set sink0 [new Agent/TCPSink]
$ns attach-agent $n5 $sink0
$ns connect $tcp0 $sink0
$tcp0 set window_ 30000
$tcp0 set fid_ 1
```

```
#Setup TCP connection S1-D2
set tcp1 [new Agent/TCP/Newreno]
$ns attach-agent $n0 $tcp1
set sink1 [new Agent/TCPSink]
$ns attach-agent $n6 $sink1
$ns connect $tcp1 $sink1
$tcp1 set window_ 30000
$tcp1 set fid_ 2
```

```
#Setup TCP connection S2-D1
set tcp2 [new Agent/TCP/Newreno]
$ns attach-agent $n1 $tcp2
set sink2 [new Agent/TCPSink]
$ns attach-agent $n5 $sink2
$ns connect $tcp2 $sink2
$tcp2 set window_ 30000
$tcp2 set fid_ 3
```

```
#Setup TCP connection S2-D2
set tcp3 [new Agent/TCP/Newreno]
$ns attach-agent $n1 $tcp3
set sink3 [new Agent/TCPSink]
$ns attach-agent $n6 $sink3
$ns connect $tcp3 $sink3
$tcp3 set window_ 30000
$tcp3 set fid_ 4
```

```
#Setup a FTP over TCP connection
set ftp0 [new Application/FTP]
$ftp0 attach-agent $tcp0
$ftp0 set type_ FTP
$ftp0 set packet_size_ 1500
```

```
#Setup a FTP over TCP connection
set ftp1 [new Application/FTP]
$ftp1 attach-agent $tcp1
$ftp1 set type_ FTP
$ftp1 set packet_size_ 1500
```

```
#Setup a FTP over TCP connection
set ftp2 [new Application/FTP]
$ftp2 attach-agent $tcp2
$ftp2 set type_ FTP
$ftp2 set packet_size_ 1500
```



```

#Setup a FTP over TCP connection
set ftp3 [new Application/FTP]
$ftp3 attach-agent $tcp3
$ftp3 set type_ FTP
$ftp3 set packet_size_ 1500

proc plotWindow {tcpSource outfile} {
    global ns

    set now [$ns now]
    set cwnd [$tcpSource set cwnd_]

    ###Print TIME CWND for gnuplot to plot progressing on CWND
    puts $outfile "$now $cwnd"

    $ns at [expr $now+0.1] "plotWindow $tcpSource $outfile"
}

set rng [new RNG]
$rng seed 0

# create a random variable that follows the uniform distribution
set loss_random_variable [new RandomVariable/Uniform]
$loss_random_variable use-rng $rng
$loss_random_variable set min_ 0
# the range of the random variable;
$loss_random_variable set max_ 100
set loss_module [new ErrorModel]
# create the error model;
$loss_module drop-target [new Agent/Null]
#a null agent where the dropped packets go to
$loss_module set rate_ 0
# right now rate is 0
# error rate will then be (0.1 = 10 / (100 - 0));
$loss_module ranvar $loss_random_variable
# attach the random variable to loss module;
$ns lossmodel $loss_module $n2 $n3

#Schedule events for the CBR and FTP agents
$ns at 0.0 "plotWindow $tcp0 $outfile1"
$ns at 0.0 "plotWindow $tcp1 $outfile2"
$ns at 0.0 "plotWindow $tcp2 $outfile3"
$ns at 0.0 "plotWindow $tcp3 $outfile4"
$ns at 1.0 "$ftp0 start"
$ns at 1.0 "$ftp1 start"
$ns at 1.0 "$ftp2 start"
$ns at 1.0 "$ftp3 start"
$ns at 99.0 "$ftp0 stop"
$ns at 99.0 "$ftp1 stop"
$ns at 99.0 "$ftp2 stop"
$ns at 99.0 "$ftp3 stop"

#Call the finish procedure after 5 seconds of simulation time
$ns at 100.0 "finish"

#Run the simulation
$ns run

```


| | | | |
|--------------------|----------|--------------------|------------|
| idletime max: | 708.8 ms | idletime max: | 708.9 ms |
| throughput: | 14 Bps | throughput: | 345248 Bps |
| RTT samples: | 3 | RTT samples: | 14596 |
| RTT min: | 30.6 ms | RTT min: | 49.1 ms |
| RTT max: | 31.1 ms | RTT max: | 212.6 ms |
| RTT avg: | 30.8 ms | RTT avg: | 49.9 ms |
| RTT stdev: | 0.3 ms | RTT stdev: | 3.6 ms |
| RTT from 3WHS: | 31.1 ms | RTT from 3WHS: | 49.8 ms |
| RTT full_sz smpls: | 2 | RTT full_sz smpls: | 1 |
| RTT full_sz min: | 30.6 ms | RTT full_sz min: | 49.8 ms |
| RTT full_sz max: | 31.1 ms | RTT full_sz max: | 49.8 ms |
| RTT full_sz avg: | 30.9 ms | RTT full_sz avg: | 49.8 ms |
| RTT full_sz stdev: | 0.0 ms | RTT full_sz stdev: | 0.0 ms |
| post-loss acks: | 0 | post-loss acks: | 52 |
| segs cum acked: | 1 | segs cum acked: | 16086 |
| duplicate acks: | 2 | duplicate acks: | 1435 |
| triple dupacks: | 0 | triple dupacks: | 52 |
| max # retrans: | 0 | max # retrans: | 1 |
| min retr time: | 0.0 ms | min retr time: | 79.9 ms |
| max retr time: | 0.0 ms | max retr time: | 158.0 ms |
| avg retr time: | 0.0 ms | avg retr time: | 88.1 ms |
| sdv retr time: | 0.0 ms | sdv retr time: | 22.5 ms |

Appendix C

```
#===== throughput.awk =====

BEGIN {
packetbuffer_flow1 = 0;
packetbuffer_flow2 = 0;
packetbuffer_flow3 = 0;
packetbuffer_flow4 = 0;
}
#body
{
        event = $1
        time = $3
        Src_node = $5
        Dest_node = $7
        Pkt_type = $9
        Size = $11
        Flow_id = $13

#===== CALCULATE throughput=====
#FLOW 1 4-5 S1-D1
if (( event == "r" ) && ( Pkt_type == "tcp" ) && ( Src_node == "4" ) && ( Dest_node ==
"5" ) && (Flow_id == "1" ) )
{
flow1_recv++;
packetbuffer_flow1 = packetbuffer_flow1 + Size;
}
#FLOW 2 4-6 S1-D2
if (( event == "r" ) && ( Pkt_type == "tcp" ) && ( Src_node == "4" ) && ( Dest_node ==
"6" ) && (Flow_id == "2" ) )
{
flow2_recv++;
packetbuffer_flow2 = packetbuffer_flow2 + Size;
}
#FLOW 3 4-5 S2-D1
if (( event == "r" ) && ( Pkt_type == "tcp" ) && ( Src_node == "4" ) && ( Dest_node ==
"5" ) && (Flow_id == "3" ) )
{
flow3_recv++;
packetbuffer_flow3 = packetbuffer_flow3 + Size;
}
#FLOW 4 4-6 S2-D2
if (( event == "r" ) && ( Pkt_type == "tcp" ) && ( Src_node == "4" ) && ( Dest_node ==
"6" ) && (Flow_id == "4" ) )
{
flow4_recv++;
packetbuffer_flow4 = packetbuffer_flow4 + Size;
}
} #body

END {
print (packetbuffer_flow1 * 8.0)/((time-1)*10^6), (packetbuffer_flow2 *
8.0)/((time-1)*10^6), (packetbuffer_flow3 * 8.0)/((time-1)*10^6), (packetbuffer_flow4
* 8.0)/((time-1)*10^6);
}
#===== Ends =====
```

Appendix D

```
#!/bin/sh
# Replace with
# $rng seed RRPP
# $loss_module set rate_ REPLACEHERE
#
# Script to automate the ECE4607 project
#
echo "Starting Simulations"

echo "NewReno TCP"

for i in 0 0.1 0.2 0.3 0.4 0.5 0.6 0.7 0.8 0.9 1.0
do
    echo "Rate = "$i
    sed 's/REPLACEHERE/'$i'/g' NewReno.tcl > NewReno_mod.tcl
    for j in 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19
    do
        sed 's/RRPP/'$j'/g' NewReno_mod.tcl
        ns NewReno_mod.tcl
        awk -f throughput.awk projout_newreno.nam >> out_newreno.txt
    done
    echo "" >> out_newreno.txt
done

echo "Reno TCP"

for i in 0 0.1 0.2 0.3 0.4 0.5 0.6 0.7 0.8 0.9 1.0
do
    echo "Rate = "$i
    sed 's/REPLACEHERE/'$i'/g' Reno.tcl > Reno_mod.tcl
    for j in 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19
    do
        sed 's/RRPP/'$j'/g' Reno_mod.tcl
        ns Reno_mod.tcl
        awk -f throughput.awk projout_reno.nam >> out_reno.txt
    done
    echo "" >> out_reno.txt
done

echo "Sack TCP"

for i in 0 0.1 0.2 0.3 0.4 0.5 0.6 0.7 0.8 0.9 1.0
do
    echo "Rate = "$i
    sed 's/REPLACEHERE/'$i'/g' Sack.tcl > Sack_mod.tcl
    for j in 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19
    do
        sed 's/RRPP/'$j'/g' Sack_mod.tcl
        ns Sack_mod.tcl
        awk -f throughput.awk projout_sack.nam >> out_sack.txt
    done
    echo "" >> out_sack.txt
done

echo "Tahoe TCP"

for i in 0 0.1 0.2 0.3 0.4 0.5 0.6 0.7 0.8 0.9 1.0
do
```

```
echo "Rate = "$i
sed 's/REPLACEHERE/'$i'/g' Tahoe.tcl > Tahoe_mod.tcl
for j in 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19
do
sed 's/RRPP/'$j'/g' Tahoe_mod.tcl
    ns Tahoe_mod.tcl
    awk -f throughput.awk projout_tahoe.nam >> out_tahoe.txt
done
echo "" >> out_tahoe.txt
done
```